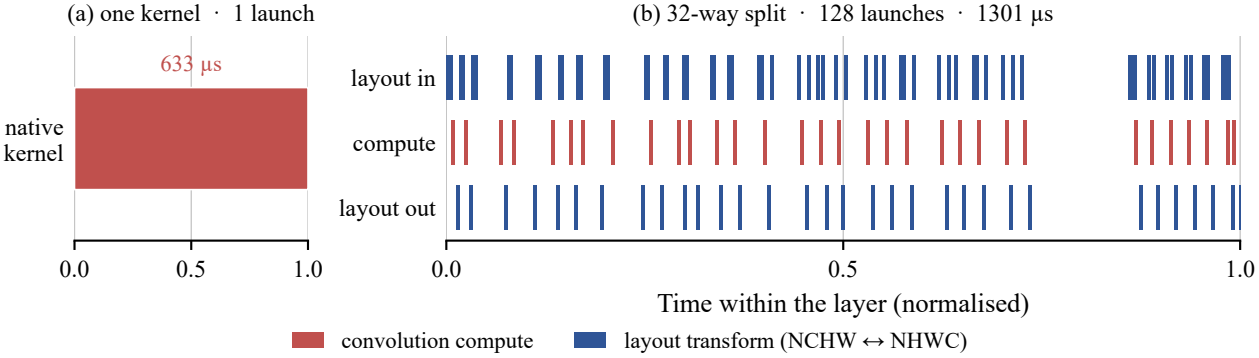
(a) one kernel · 1 launch
633 µs
native kernel
0.0
0.5
1.0
(b) 32-way split · 128 launches · 1301 µs
layout in
compute
layout out
0.0
0.5
1.0
Time within the layer (normalised)
convolution compute
layout transform (NCHW ↔ NHWC)